

WHestBench Paper A — reviewer questions

The request is intentionally bounded. A reviewer need not audit the complete experiment ledger.

Theorem and scope

Is the static estimator class stated naturally and without hidden assumptions?

Does the nonnegative result support “within 0.0233242% of optimum,” given the stated interval-verification dependency?

Is the arbitrary-signed theorem correctly separated from the stronger nonnegative theorem?

Is it correct to exclude the earlier arbitrary-total-mass corollary until its standalone witness is recovered, and should the mass-one restriction be more prominent?

Are the finite-width, adaptive, nonlinear, and network-dependent escape classes sufficiently prominent?

Proof architecture

Is the all-degree degree-five auxiliary optimization adequately justified by the Hermite contact conditions, exact quadrature, and reduced-cost sign argument?

Is the positive-index lemma and its transfer through atomic inertia correct?

Is the sign-count hierarchy a useful corollary or better placed in an appendix?

Is the Schur–Horn shared-Gram sharpness theorem stated at the right generality?

Is the equality-to-zero-code argument correct under the full-row-rank condition?

Does the consecutive-Gegenbauer common-zero obstruction establish only strict nonattainment, as claimed?

Is the coalescing-cancellation example sufficient to explain why strictness does not give a uniform unrestricted signed gap?

Computer-assisted evidence

Which numerical interval claims must be independently reconstructed before a preprint?

Is exact-rational replay of the witness sufficient once the kernel intervals are independently certified?

What should be placed in the main paper versus a verification appendix or repository?

Positioning

Is the contribution best positioned as spherical cubature, kernel quadrature, design theory, or a neural-kernel application?

Which related work or prior near-optimality theorem is missing?

Which venue or workshop is the best initial target?

Artifact and wording checks

Does the frozen-witness inertia strengthening correctly convert the released comparison allocation from rank N to positive index $N-1$?

Is the distinction between a factor cap 1.067167... and a 6.293983...% Kerdock-risk reduction stated correctly?

Does the endpoint certificate adequately support equality only at the three Hermite contact points?

Is the distinction from exact SDP packing optimality of related Kerdock spherical codes clear enough?